



*Figure 8.* Learning curves of MaxEntDP and SAC on the dog and humanoid benchmarks from DeepMind Control suite. The solid lines are the means and the shaded region represents the standard deviations over five runs.